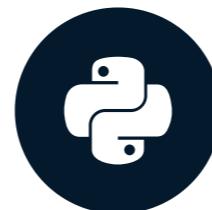


Processing and classifying images

WORKING WITH HUGGING FACE



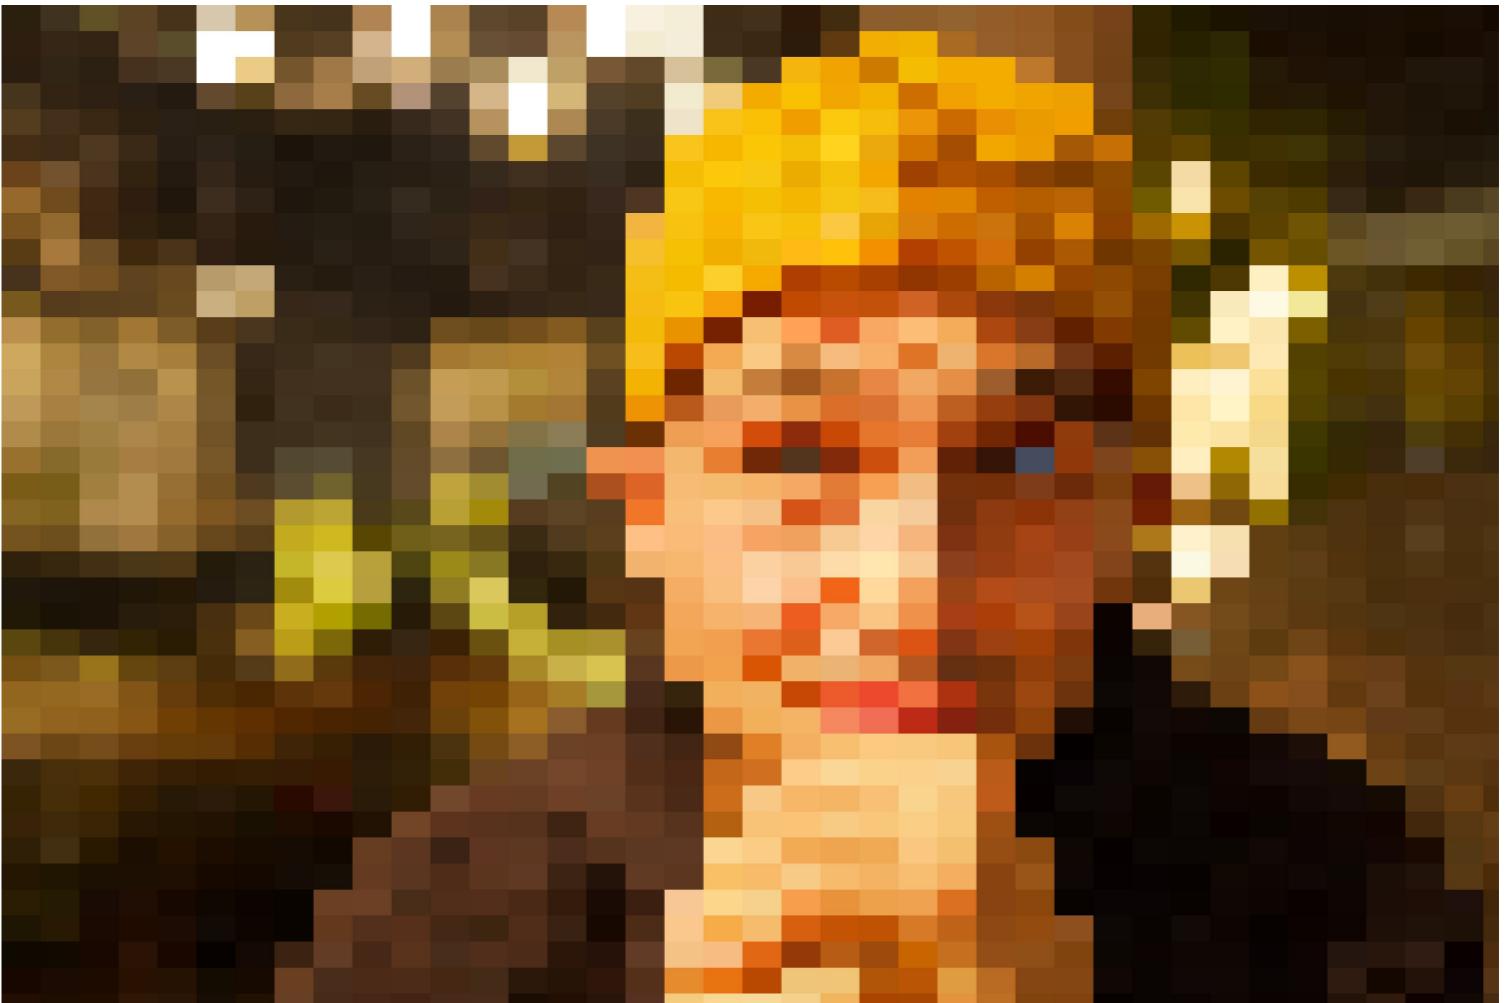
Jacob H. Marquez
Lead Data Engineer

Pixels in images



Pixels in images

- Images are made of pixels
- Pixels contain info about color and more
- Number of pixels from 1 to infinite



Pixels in images

- Images are made of pixels
- Pixels contain info about color and more
- Number of pixels from 1 to infinite
- More pixels = higher resolution



Importance of pre-processing images

- Models have expectations for the images
- Maintains consistency
- Target specific parts of images

Common pre-processing steps

- Cropping
- Resizing

Common pre-processing steps

Cropping



Common pre-processing steps

Cropping

- Removing unwanted portions of the image

Resizing

- Change dimensions of the image



Common pre-processing steps

Cropping

- Removing unwanted portions of the image

Resizing

- Change dimensions of the image

Transforming images

```
from transformers import ImageTransforms
```

```
from PIL import Image
```

```
original_image = Image.open("my_image.jpeg")
```

```
import numpy as np
```

```
image_array = np.array(original_image)
```

```
array([[[43, 32, 14], ..., [77, 47, 23]]], dtype=uint8)
```

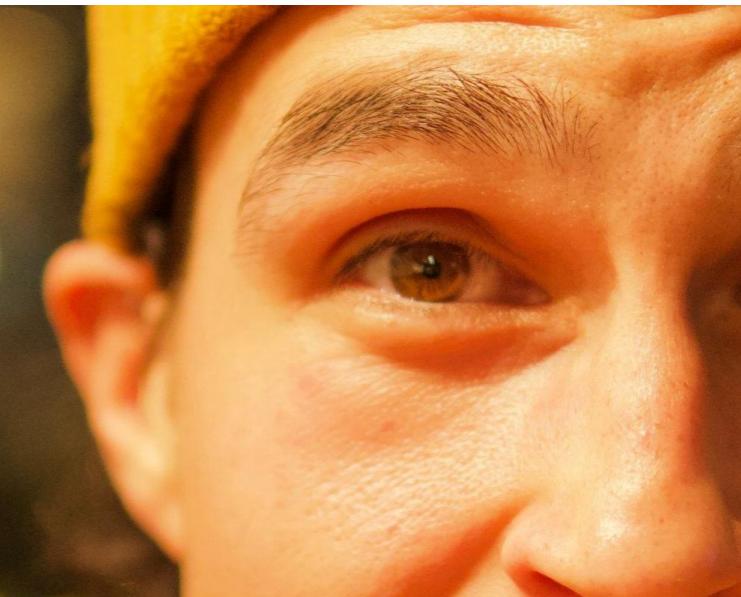
Cropping an image

```
cropped_image = image_transforms \  
    .center_crop(  
        image=np.ndarray,  
        size=(1000, 1000)  
)
```

Original image

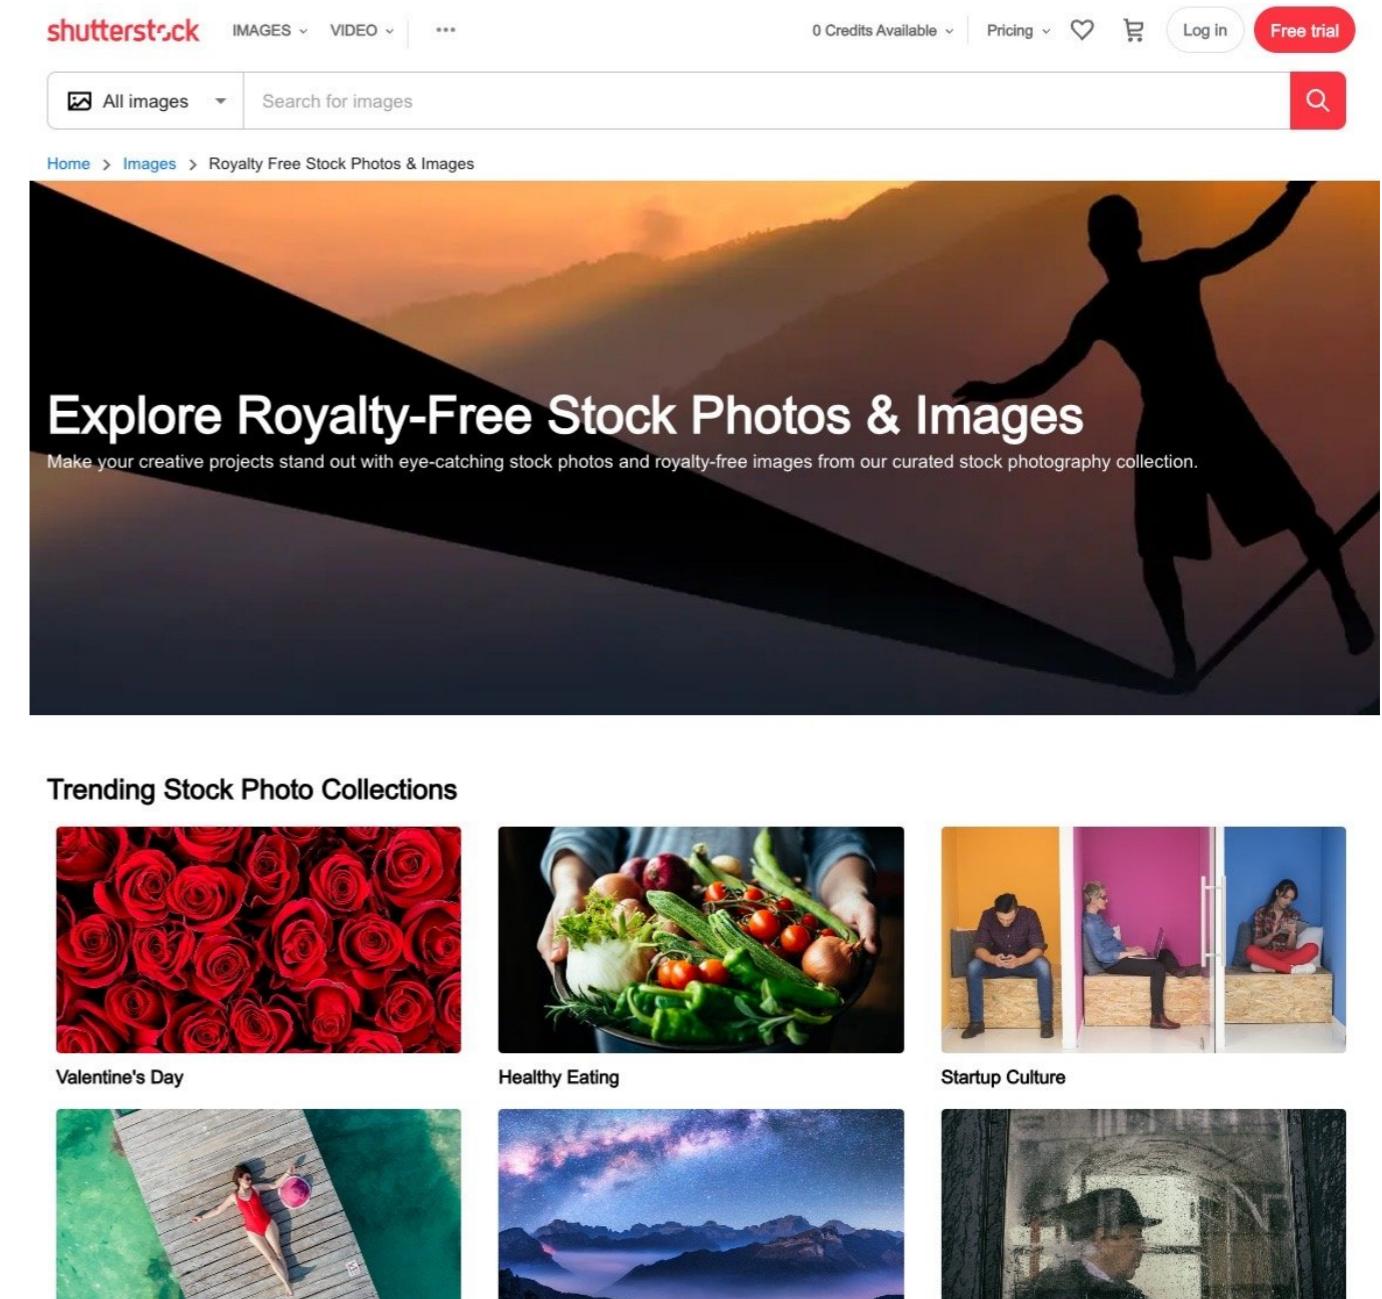


Cropped image



What is image classification?

- The process of assigning labels to an image
- Stock photography



¹ <https://www.shutterstock.com/photos>

What is image classification?

- The process of assigning labels to an image
- Stock photograph
- Agriculture



¹ <https://phys.org/news/2021-03-crop-diseasesthere-app.html>

What is image classification?

- The process of assigning labels to an image
- Stock photograph
- Agriculture
- Wildlife monitoring



Image classification pipeline

```
from transformers import pipeline  
  
classifier = pipeline(task="image-classification",  
                      model="google/vit-base-patch16-224")
```

Image classification pipeline

```
# http link to image  
classifier("https://www.pexels.com/photo/elephants.jpeg")
```

```
# local path to image  
classifier("/my_folder/elephants.jpeg")
```

```
# PIL image  
from PIL import Image  
image = Image.open("/my_folder/elephants.jpeg")  
classifier(image)
```

```
classifier(["image1.jpeg", "image2.jpeg"])
```

Top K parameter

```
from transformers import pipeline

classifier = pipeline(task="image-classification",
                      model="google/vit-base-patch16-224")

results = classifier(image, top_k=2)

print(results[0]['label'])
```

African elephant, Loxodonta africana

Let's practice!

WORKING WITH HUGGING FACE

Question answering and multi-modal tasks

WORKING WITH HUGGING FACE



Jacob H. Marquez
Lead Data Engineer

What is document question and answering?

- Answering questions about content of document
- Document is a text-based image
- Question is specific to the document
- Answer can be direct quote or paraphrased response

Memo to: Distribution
July 19, 2002
Page 2

Action Steps:

- Please interpret this directive in its broadest sense to prevent the deletion or destruction of any recorded information and data relating in any way to Actos.
- Please take steps immediately to preserve such documents and data within your department.
- Please distribute this memo to members of your group and advise them of the importance of following these instructions.

* * * *

If you have any questions regarding the implementation of this directive, please contact me.

Question: "What are the action steps?"

Visual question and answering

- Ask a question of an actual image or video
- Pipeline requires the image and question



Question: "What type of animal is in this picture?"

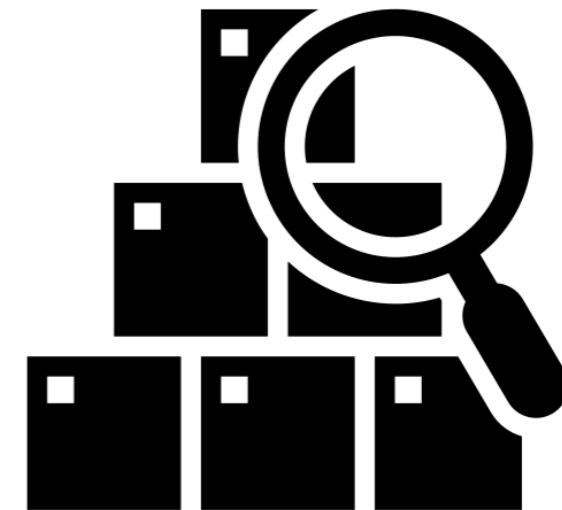
Use cases: information retrieval



Customer
Support



Legal Compliance



Search

Preprocessing for multi-modal tasks

- Multi-modal
- Process each data type
- Tokenization for text
- Resizing for images

Memo to: Distribution
July 19, 2002
Page 2

Action Steps:

- Please interpret this directive in its broadest sense to prevent the deletion or destruction of any recorded information and data relating in any way to Actos.
- Please take steps immediately to preserve such documents and data within your department.
- Please distribute this memo to members of your group and advise them of the importance of following these instructions.

If you have any questions regarding the implementation of this directive, please contact me.



Document Q&A pipeline

```
from transformers import pipeline

dqa = pipeline(
    task="document-question-answering",
    model="naver-clova-ix/donut-base-finetuned-docvqa")
document_image = "memo.jpg"
question_text = "What is this memo about?"

results = dqa(document_image, question_text)
```

Results of document Q&A pipeline

```
print(results)
```

```
{  
    "score": 0.789,  
    "start": 1,  
    "end": 2,  
    "answer": "distribution",  
    "words": [102]  
}
```

```
dqa(image=image,  
     question=question,  
     max_answer_len=15)
```

- `score` is the probability of the answer
 - `answer` is the answer to the question
 - `start` is the start word index of the answer
 - `end` is the last word index of the answer
 - `words` is a list of all indices for each word in the answer
-
- 79% probability the memo is about distribution

Visual Q&A pipeline

```
from transformers import pipeline  
  
vqa = pipeline(  
    task="visual-question-answering",  
    model="dandelin/vilt-b32-finetuned-vqa"  
)
```

```
result = vqa(  
    image="image.jpeg",  
    question="what's the person wearing?")
```



Results of visual Q&A pipeline

```
print(result)
```

```
[  
    {'score': 0.9795706272125244,  
     'answer': 'hat'  
    },  
    ...  
    {'score': 0.02153933234512806,  
     'answer': 'hoodie'  
    }  
]
```

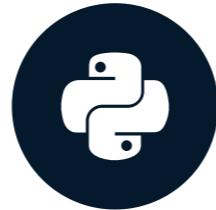
- **label** label identified by the model
- **score** probability of the label from the model

Let's practice!

WORKING WITH HUGGING FACE

Audio classification

WORKING WITH HUGGING FACE



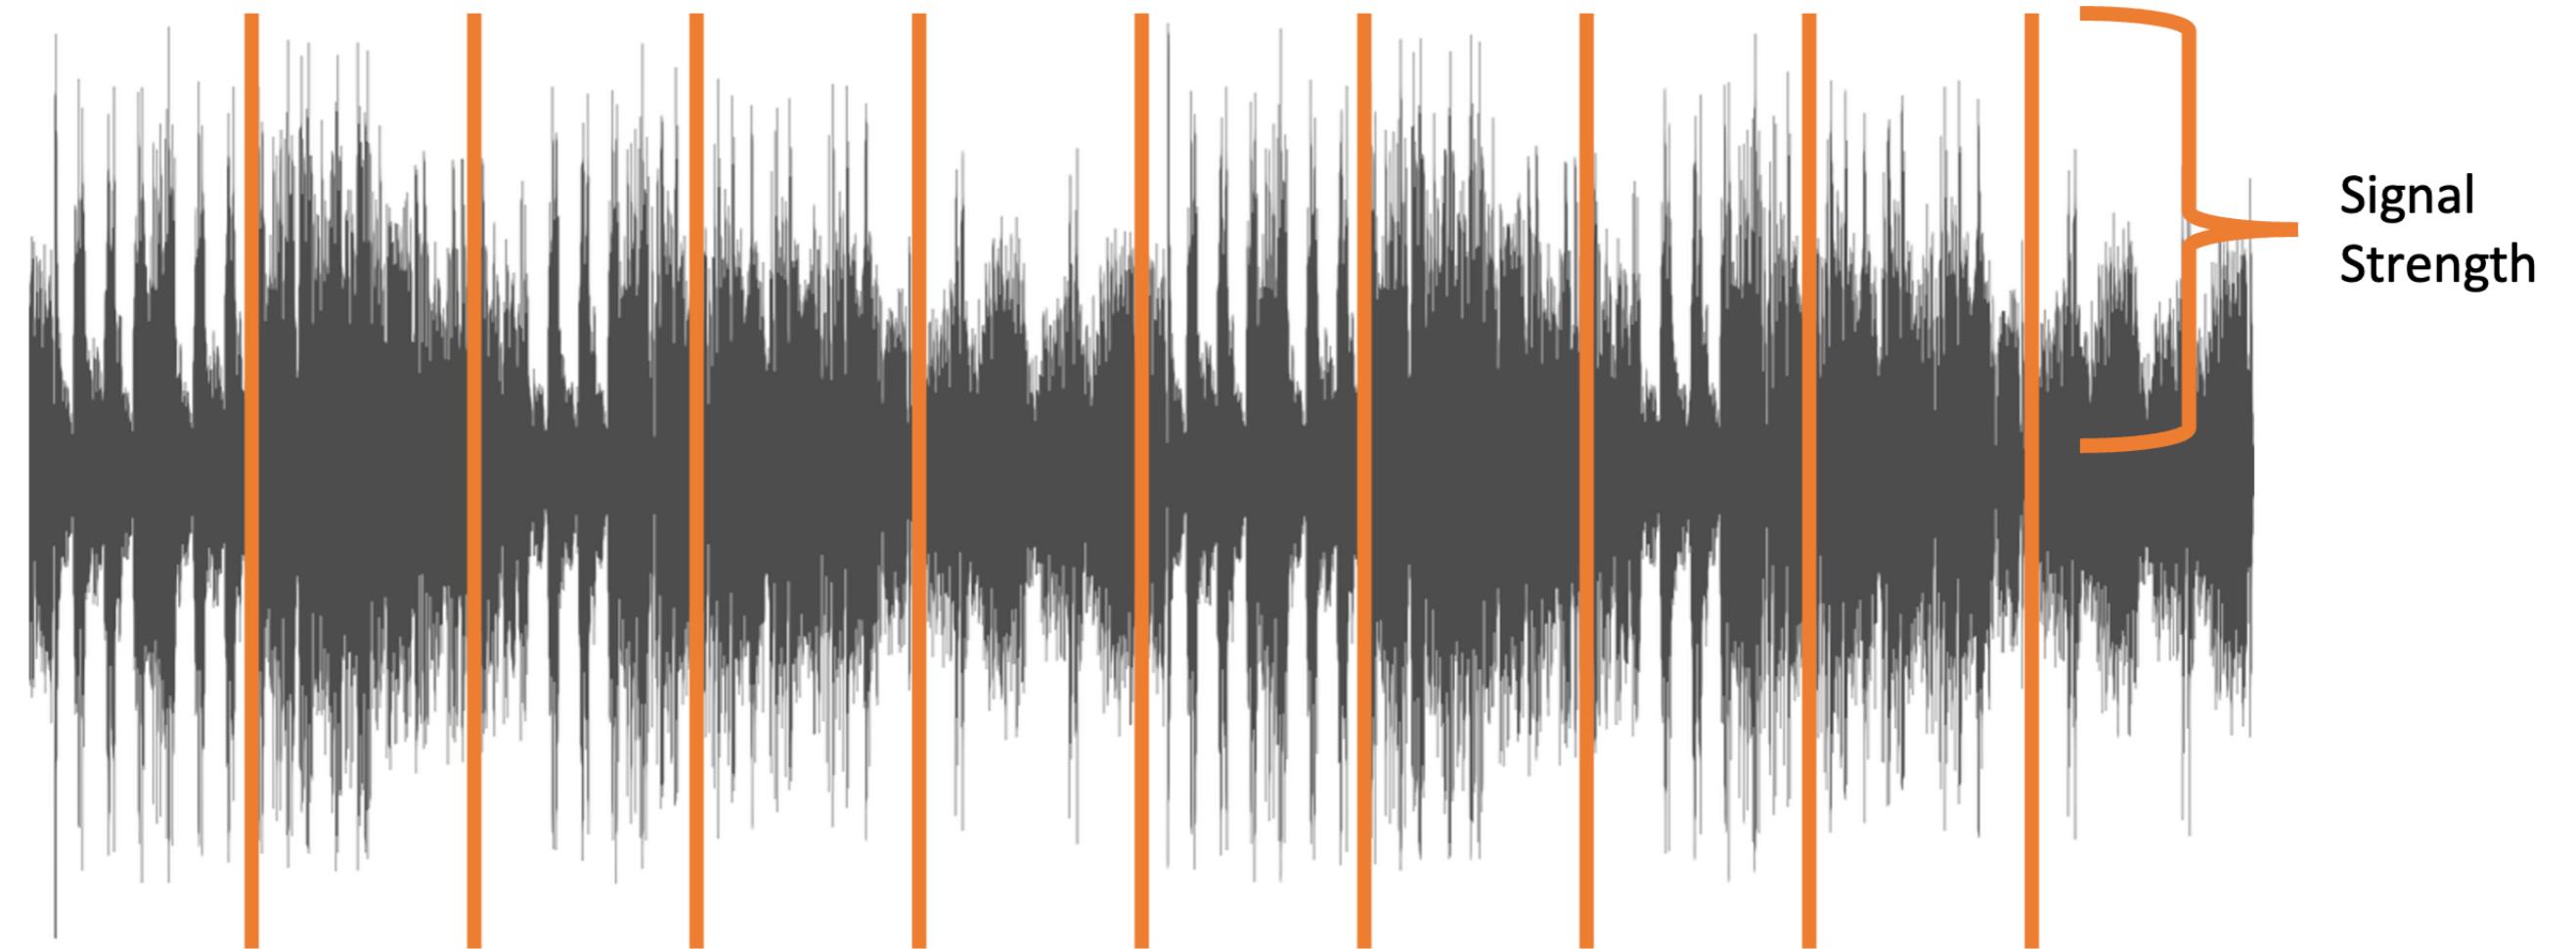
Jacob H. Marquez
Lead Data Engineer

What is audio data?

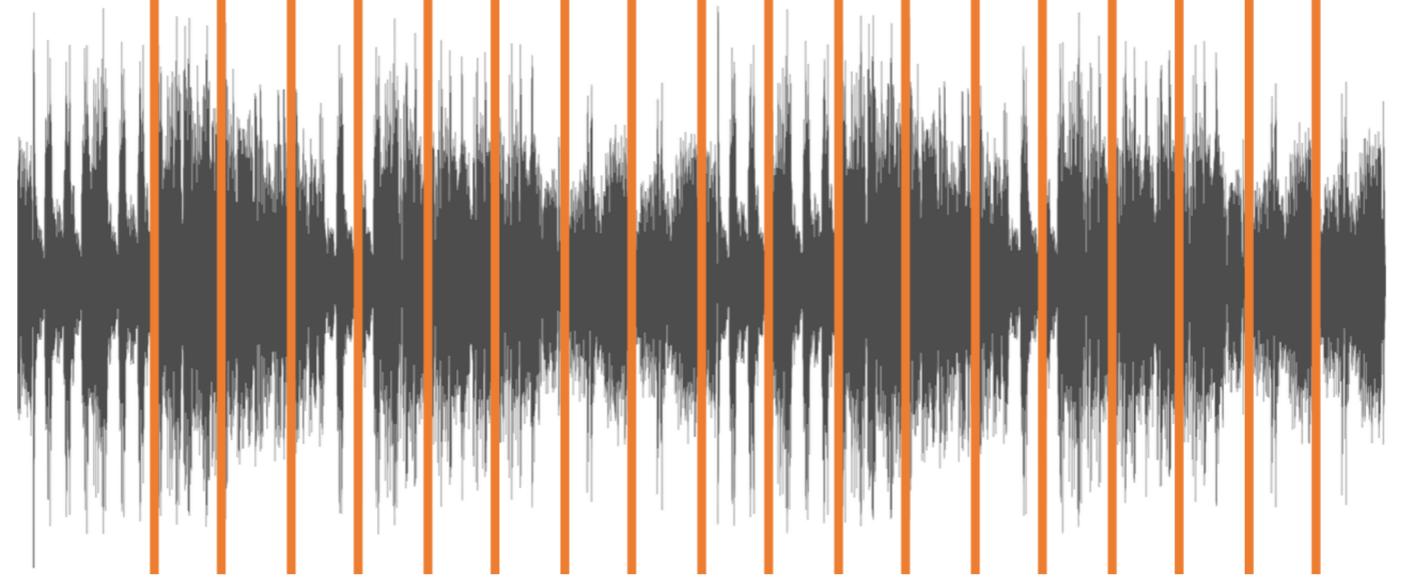
- Continuous signal called a sound wave
- Length and amplitude
- Converted into series of discrete values
- Results in a digital representation
- Conversion helps ML algorithms process and use audio
- Sampling is an important step



The importance of sampling

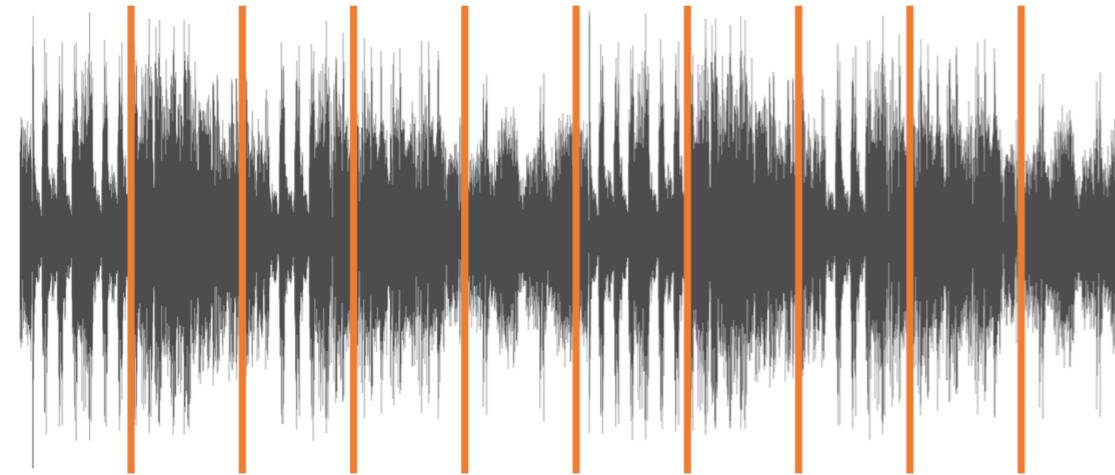


The importance of sampling

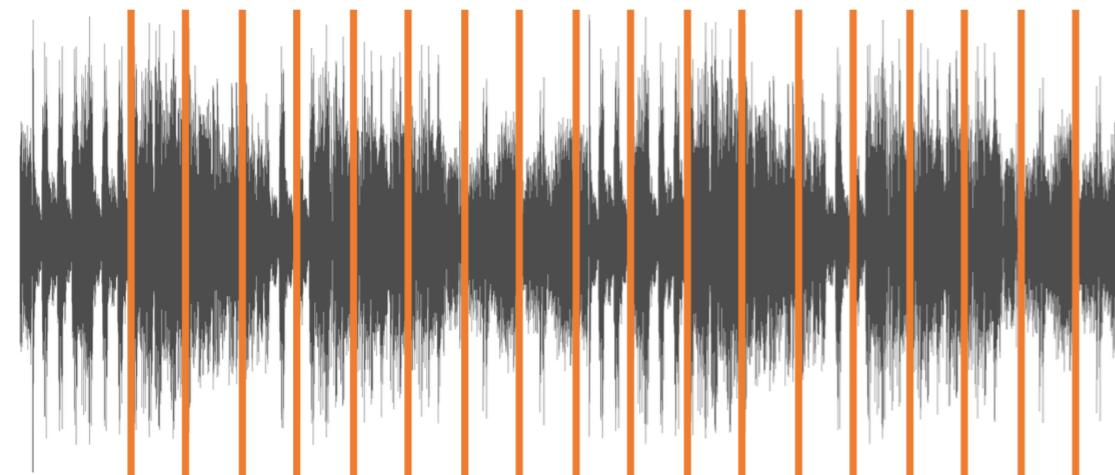


- Speech models trained at 16kHz
- Sampling rate specified in the model card

Resampling



- Aligns sample rates across all files
- Ensures consistency
- Helps with processing



Resampling using Hugging Face

```
from datasets import Audio  
  
songs = songs.cast_column("audio", Audio(sampling_rate=16_000))
```

Finding the sampling rate:

```
print(songs[0]["audio"]["sampling_rate"])
```

```
16_000
```

Filtering

Benefits

- Ensure enough audio for inference or training
- Reduce computation by limiting file size

```
import librosa
```

Filtering

```
durations = []

for row in songs["path"]:
    durations.append(librosa.get_duration(path=row))

songs.add_column("duration", durations)

songs = dataset.filter(
    lambda d: d < 10.0, input_columns=["duration"]
)
```

What is audio classification?

Definition: process of assigning one or more labels to audio clips based on its content



Language
Identification

What is audio classification?

Definition: process of assigning one or more labels to audio clips based on its content



Language
Identification



Environmental
Sounds

What is audio classification?

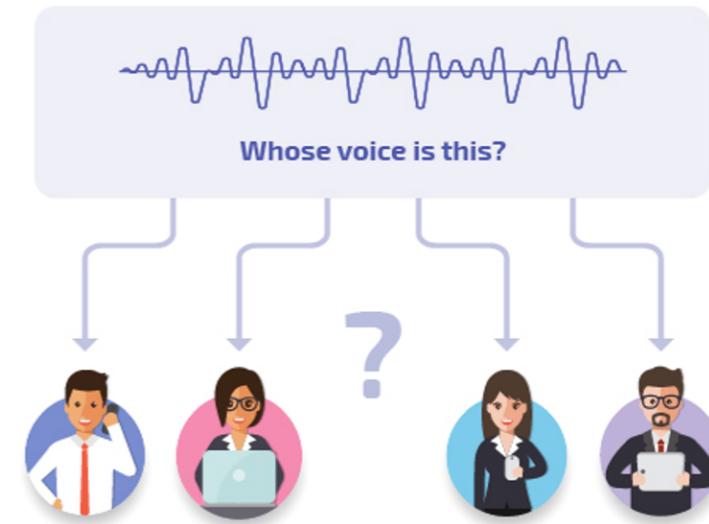
Definition: process of assigning one or more labels to audio clips based on its content



Language
Identification



Environmental
Sounds



Speaker
Identification

Using Hugging Face pipelines

```
from transformers import pipeline

classifier = pipeline(task="audio-classification",
                      model="superb/wav2vec2-base-superb-ks")

genreClassifier = pipeline(task="audio-classification",
                           model="mtg-upf/discogs-maest-30s-pw-73e-ts")
```

Using Hugging Face pipelines

```
audio = songs[0]['audio']['array']
prediction = genreClassifier(audio)

print(prediction)
```

```
[
{'score': 0.07648807018995285, 'label': 'Non-Music---Field Recording'},
...
{'score': 0.05880315974354744, 'label': 'Electronic---Noise'}
]
```

¹ <https://huggingface.co/mtg-upf/discogs-maest-30s-pw-73e-ts>

Using Hugging Face pipelines

```
prediction = genreClassifier(audio, top_k=1)
```

```
[  
{'score': 0.07648807018995285, 'label': 'Non-Music---Field Recording'}  
]
```

Let's practice!

WORKING WITH HUGGING FACE

Automatic speech recognition

WORKING WITH HUGGING FACE



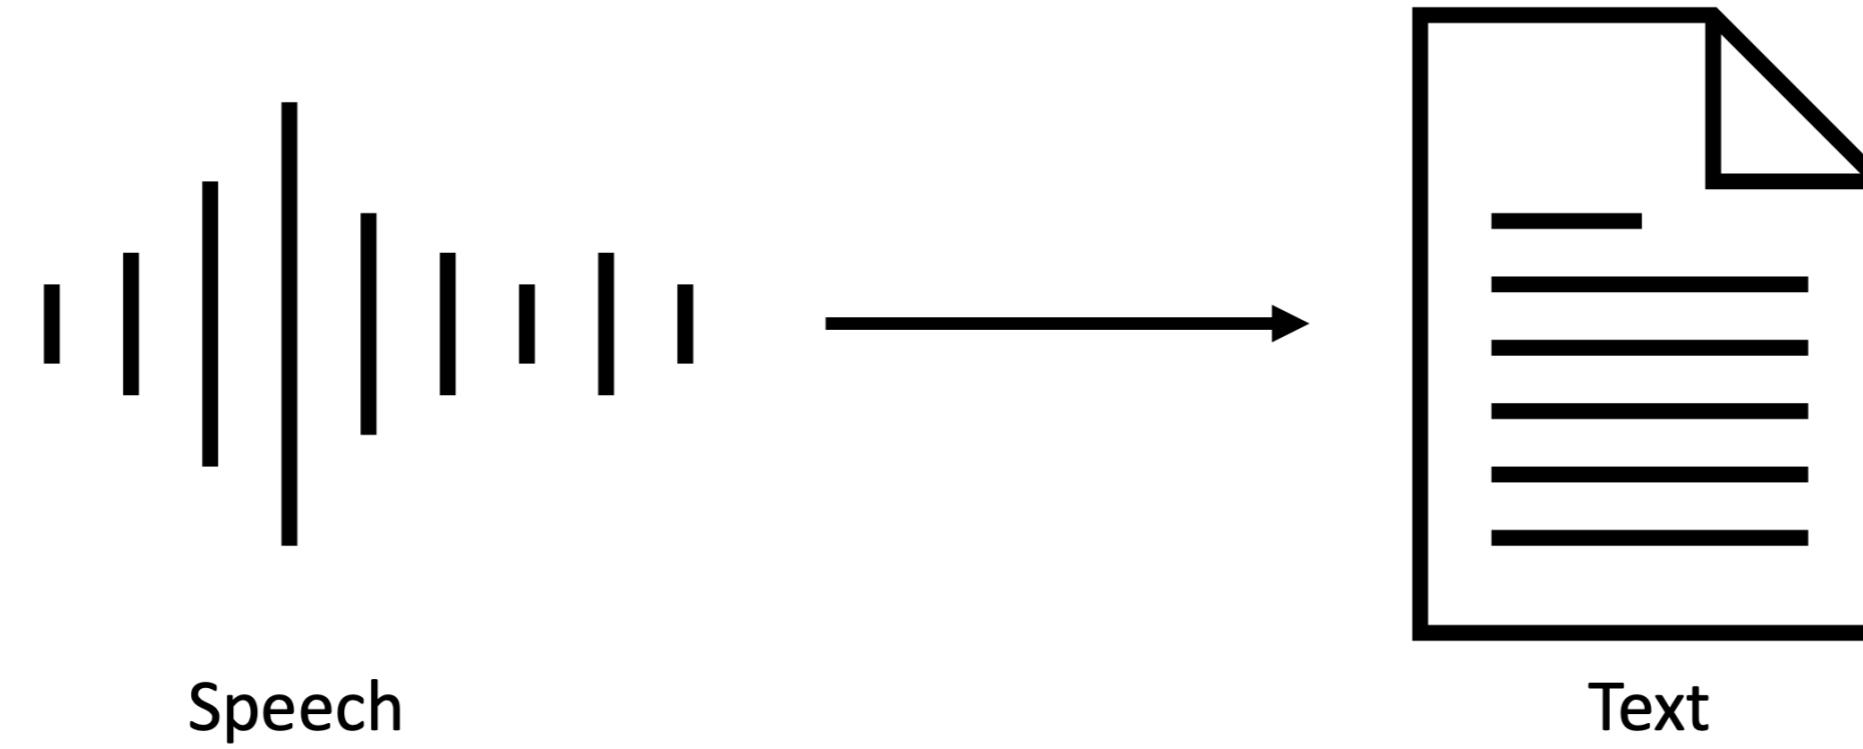
Jacob H. Marquez
Lead Data Engineer

What is automatic speech recognition?



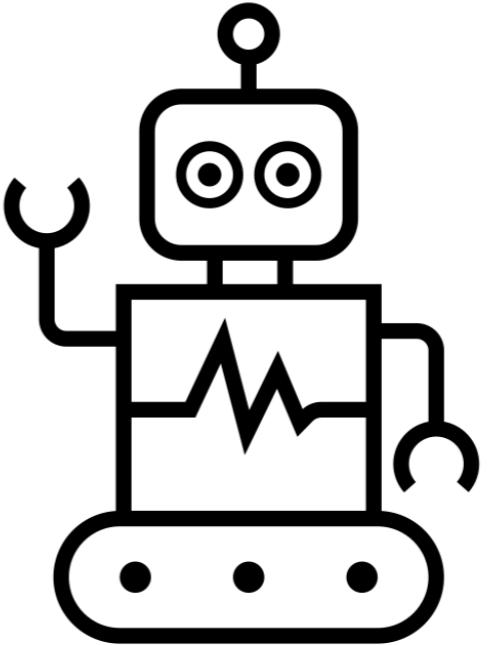
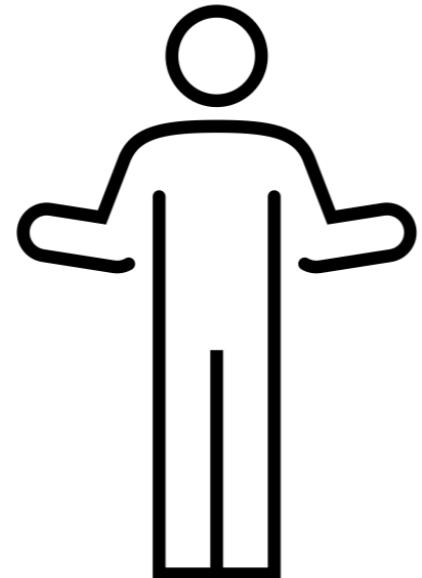
Speech

What is automatic speech recognition?



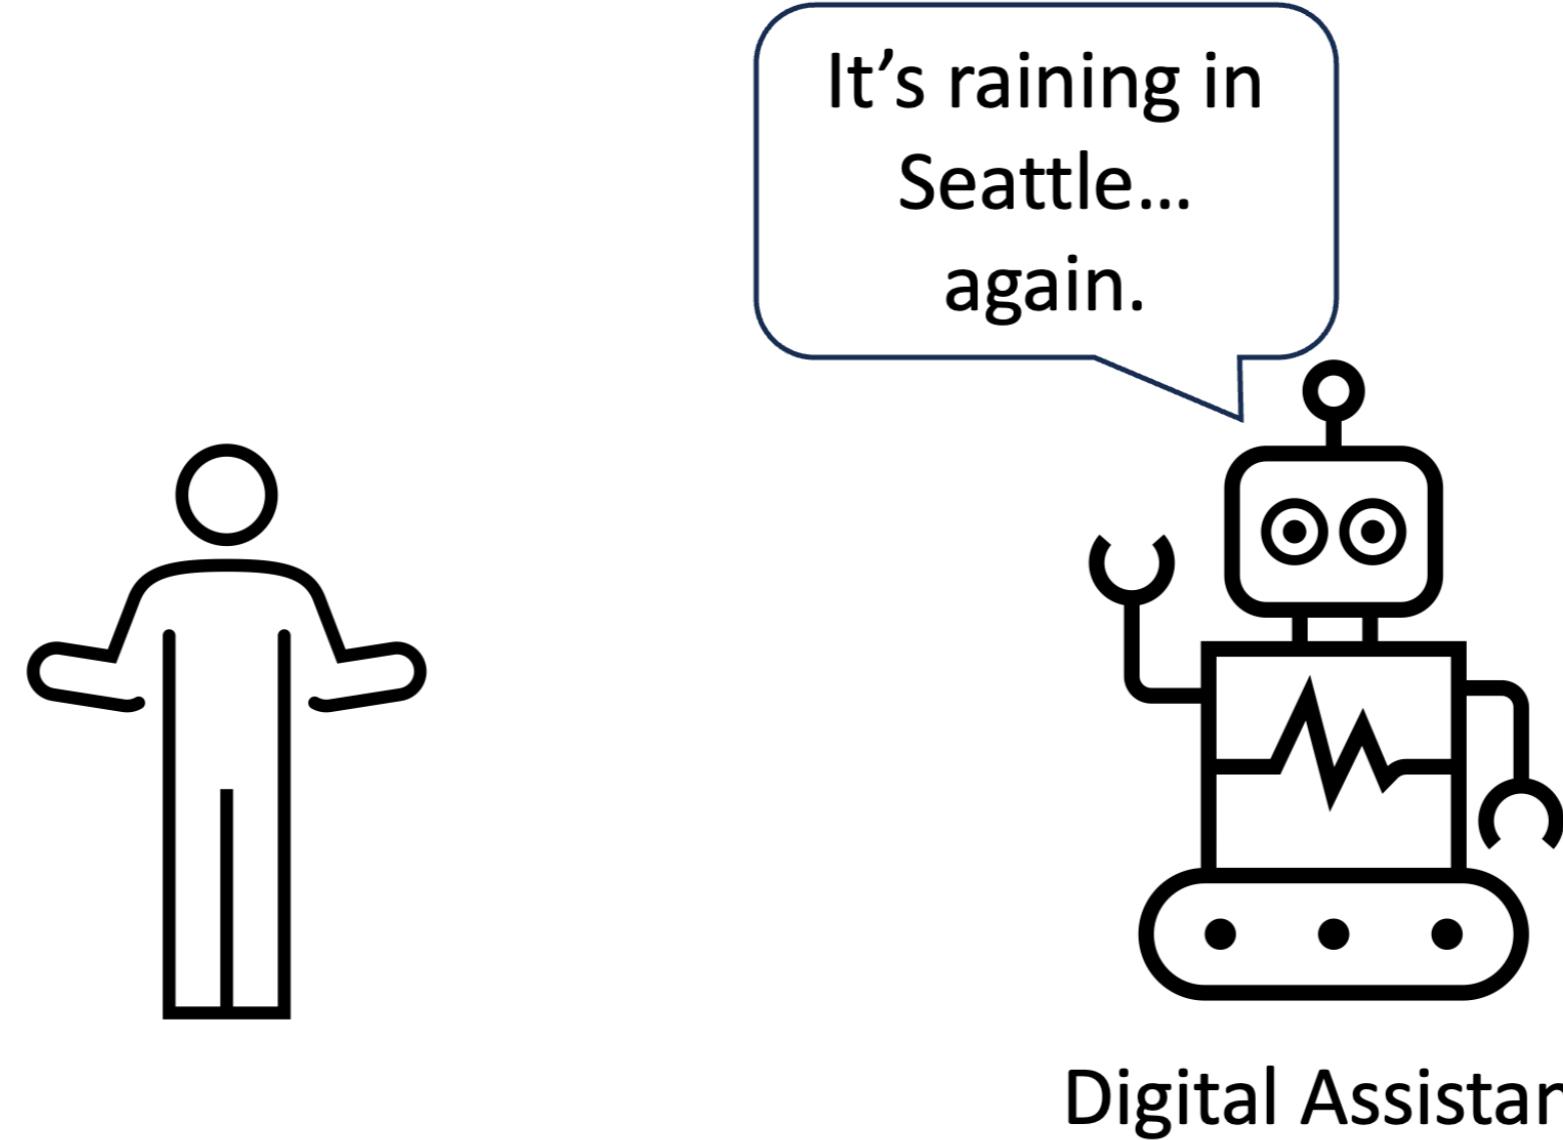
Use cases of ASR

What's the
weather
today?



Digital Assistant

Use cases of ASR



Use cases for ASR



- Create transcripts
- Finding relevant documentation and solutions

Use cases for ASR

Assumptions about the SOLUTION



Problem vs Feature focus Find real friction points by going to prod The ML Twist

solutions need to solve for the ml twist
and we can't just say ah

Models for ASR



Wav2Vec

```
model_id:  
"facebook/wav2vec2-base-960h"
```

Models for ASR



Wav2Vec

```
model_id:  
"facebook/wav2vec2-base-960h"
```



Whisper

```
model_id:  
"openai/whisper-base"
```

- Whisper performs better with punctuation and casing.

Instantiating a pipeline for ASR

```
transcriber = pipeline(task="automatic-speech-recognition",
                      model="facebook/wav2vec2-base-960h")

# Path to audio file
transcriber("my_audio.wav")

# Numpy array
transcriber(numpy_audio_array)

# Dictionary
transcriber({"sampling_rate" = 16_000, "raw" = "my_audio.wav"})
```

Results from a pipeline

```
sampling_rate = 16_000  
dataset = dataset.cast_column("audio", Audio(sampling_rate=sampling_rate))
```

```
input = data[0]['audio']['array']
```

```
prediction = transcriber(input)
```

```
print(prediction)
```

```
"what game do you want to play"
```

Predicting over a dataset

```
def data():
    for i in range(dataset):
        yield dataset[i]['audio']['array'], dataset[i]['sentence'].lower()

output = []

for audio, sentence in data():
    prediction = transcriber(audio)
    output.append((prediction, sentence))
```

```
[("what a nice black shirt", "what a nice blue shirt"), ...]
```

Evaluating ASR systems

- Word Error Rate (WER)
- Based on Levenshtein Distance
- Metric for the difference between two sequences

$$\frac{(Substitutions + Insertions + Deletions)}{Number of Words Spoken}$$

- Range from 0 to 1
- Smaller value indicates closer similarity

¹ https://en.wikipedia.org/wiki/Levenshtein_distance

Word Error Rate

Correct

I love DataCamp
courses on AI!

Prediction

I love DataCamp
portraits on **hay**!

- 2 substitutions required to match to correct

$$2 / 6 = 0.33$$

Computing WER using Hugging Face

```
from evaluate import load

# Instantiate word error rate metric
wer = load("wer")

# Save true sentence as reference
reference = data[0]['sentence']

predictions = "I love DataCamp portraits on hay"
```

¹ <https://huggingface.co/spaces/evaluate-metric/wer>

Computing WER using Hugging Face

```
# Compute the WER between predictions and reference  
wer_score = wer.compute(  
    predictions=[prediction],  
    references=[reference]  
)  
  
print(wer_score)
```

0.33

Let's practice!

WORKING WITH HUGGING FACE